
ARTIFICIAL NEURAL NETWORKS

Classification Using Artificial Neural Networks with Hyperparameter Tuning on Alphabets Data

Overview

In this assignment, you will be tasked with developing a classification model using Artificial Neural Networks (ANNs) to classify data points from the "Alphabets_data.csv" dataset into predefined categories of alphabets. This exercise aims to deepen your understanding of ANNs and the significant role hyperparameter tuning plays in enhancing model performance.

Dataset: "Alphabets_data.csv"

The dataset provided, "Alphabets_data.csv", consists of labeled data suitable for a classification task aimed at identifying different alphabets. Before using this data in your model, you'll need to preprocess it to ensure optimal performance.

Tasks

1. Data Exploration and Preprocessing

- Begin by loading and exploring the "Alphabets_data.csv" dataset. Summarize its key features such as the number of samples, features, and classes.
- Execute necessary data preprocessing steps including data normalization, managing missing values.

Answer :

Code used :

```
import pandas as pd
import pprint
import os

# --- File paths ---
input_path = r"D:\DATA SCIENCE\ASSIGNMENTS\18 neural networks\Neural networks\Alphabets_data.csv"
output_folder = r"d:\DATA SCIENCE\ASSIGNMENTS\18 neural networks\Neural networks"
output_file = os.path.join(output_folder, "dataset_overview.txt")

# --- Load dataset ---
alphabets_data_df = pd.read_csv(input_path)

# --- Dataset summary ---
dataset_overview = {
    "Shape": alphabets_data_df.shape,
    "Column_Names": alphabets_data_df.columns.tolist(),
    "Missing_Values": alphabets_data_df.isnull().sum().to_dict(),
    "Data_Types": alphabets_data_df.dtypes.astype(str).to_dict(),
    "Sample_Rows": alphabets_data_df.head().to_dict(orient="records")
}

# --- Save output to file ---
with open(output_file, "w") as f:
    f.write("==== DATASET OVERVIEW ====\n\n")
    pprint.pprint(dataset_overview, stream=f, sort_dicts=False)

print(f"☑ Dataset summary saved successfully at:{output_file}")
```

==== DATASET OVERVIEW ====

```
{"Shape": (20000, 17),
'Column_Names': ['letter',
                  'xbox',
                  'ybox',
                  'width',
                  'height',
                  'onpix',
                  'xbar',
                  'ybar',
                  'x2bar',
                  'y2bar',
                  'xybar',
                  'x2ybar',
                  'xy2bar',
                  'xedge',
                  'xedgey',
                  'yedge',
                  'yedgey'],
'Missing_Values': {'letter': 0,
                   'xbox': 0,
                   'ybox': 0,
                   'width': 0,
                   'height': 0,
                   'onpix': 0,
                   'xbar': 0,
                   'ybar': 0,
                   'x2bar': 0,
                   'y2bar': 0,
                   'xybar': 0,
                   'x2ybar': 0,
                   'xy2bar': 0,
                   'xedge': 0,
                   'xedgey': 0,
                   'yedge': 0,
                   'yedgey': 0},
'Data_Types': {'letter': 'object',
                'xbox': 'int64',
                'ybox': 'int64',
                'width': 'int64',
                'height': 'int64',
                'onpix': 'int64',
                'xbar': 'int64',
                'ybar': 'int64',
                'x2bar': 'int64',
                'y2bar': 'int64',
                'xybar': 'int64',
                'x2ybar': 'int64',
                'xy2bar': 'int64',
                'xedge': 'int64',
                'xedgey': 'int64',
                'yedge': 'int64',
                'yedgey': 'int64'},
'Sample_Rows': [{"letter": "T",
```

```
'xbox': 2,
'ybox': 8,
'width': 3,
'height': 5,
'onpix': 1,
'xbar': 8,
'ybar': 13,
'x2bar': 0,
'y2bar': 6,
'xybar': 6,
'x2ybar': 10,
'xy2bar': 8,
'xedge': 0,
'xedgey': 8,
'yedge': 0,
'yedgex': 8},
{'letter': 'I',
'xbox': 5,
'ybox': 12,
'width': 3,
'height': 7,
'onpix': 2,
'xbar': 10,
'ybar': 5,
'x2bar': 5,
'y2bar': 4,
'xybar': 13,
'x2ybar': 3,
'xy2bar': 9,
'xedge': 2,
'xedgey': 8,
'yedge': 4,
'yedgex': 10},
{'letter': 'D',
'xbox': 4,
'ybox': 11,
'width': 6,
'height': 8,
'onpix': 6,
'xbar': 10,
'ybar': 6,
'x2bar': 2,
'y2bar': 6,
'xybar': 10,
'x2ybar': 3,
'xy2bar': 7,
'xedge': 3,
'xedgey': 7,
'yedge': 3,
'yedgex': 9},
{'letter': 'N',
'xbox': 7,
'ybox': 11,
'width': 6,
'height': 6,
'onpix': 3,
```

```

'xbar': 5,
'ybar': 9,
'x2bar': 4,
'y2bar': 6,
'xybar': 4,
'x2ybar': 4,
'xy2bar': 10,
'xedge': 6,
'xedgey': 10,
'yedge': 2,
'yedgex': 8},
{'letter': 'G',
'bbox': 2,
'ybox': 1,
'width': 3,
'height': 1,
'onpix': 1,
'xbar': 8,
'ybar': 6,
'x2bar': 6,
'y2bar': 6,
'xybar': 6,
'x2ybar': 5,
'xy2bar': 9,
'xedge': 1,
'xedgey': 7,
'yedge': 5,
'yedgex': 10}]}

```

2. Model Implementation

- Construct a basic ANN model using your chosen high-level neural network library. Ensure your model includes at least one hidden layer.
- Divide the dataset into training and test sets.
- Train your model on the training set and then use it to make predictions on the test set.

Answer :

Python script that:

- Loads Alphabets_data.csv (your path or the uploaded /mnt/data copy),
- Preprocesses (label encoding → one-hot, StandardScaler),
- Splits data (train/test),
- Builds a simple Keras ANN with **one hidden layer** (plus optional deeper variant),
- Trains the model with EarlyStopping,
- Saves trained model, training history (.png) and test predictions (.csv) into your folder d:/python apps/neural networks,
- Prints evaluation metrics (accuracy, precision, recall, F1) and confusion matrix.

Drop this into a .py file (e.g., ann_alphabets.py) and run it in your venv. If you prefer I can tweak architecture/hyperparams after you peek at results.

Code used :

.....

ANN for Alphabets classification (basic implementation)

- Expects Alphabets_data.csv available at input_path (change if needed)

- Saves outputs (model, metrics, plots, preds) to output_folder
- Uses TensorFlow / Keras

Requirements:

```
pip install numpy pandas scikit-learn matplotlib tensorflow
""""
```

```
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

# -----
# Paths - change if needed
# -----
# If you want to use the copy uploaded to this environment, set input_path =
#/mnt/data/Alphabets_data.csv"
input_path = r"D:\DATA SCIENCE\ASSIGNMENTS\18 neural networks\Neural
networks\Alphabets_data.csv"
# Save outputs to this folder (as you requested)
output_folder = r"d:\python apps\neural networks"

os.makedirs(output_folder, exist_ok=True)

# -----
# Load dataset
# -----
df = pd.read_csv(input_path)
print(f'Loaded data: {df.shape[0]} rows, {df.shape[1]} columns')

# -----
# Preprocessing
# -----
# 1) Separate X and y
target_col = "letter"
if target_col not in df.columns:
    raise ValueError(f'Target column '{target_col}' not found in CSV.')

X = df.drop(columns=[target_col])
y = df[target_col]

# 2) Label encode target and one-hot encode for Keras
le = LabelEncoder()
y_int = le.fit_transform(y)          # integers 0..25
num_classes = len(le.classes_)
y_cat = tf.keras.utils.to_categorical(y_int, num_classes=num_classes)

print(f'Detected {num_classes} classes: {list(le.classes_)}')
```

```

# 3) Feature scaling
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# 4) Train-test split (stratify to keep class balance)
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y_cat, test_size=0.2, random_state=42, stratify=y_int
)

print("Train shape:", X_train.shape, "Test shape:", X_test.shape)

# Optionally save the scaler and label encoder for later inference
import joblib
joblib.dump(scaler, os.path.join(output_folder, "scaler.joblib"))
joblib.dump(le, os.path.join(output_folder, "label_encoder.joblib"))

# -----
# Build model (basic: one hidden layer)
# -----
input_dim = X_train.shape[1]
hidden_units = 128 # sensible default for 16 features
dropout_rate = 0.2

def build_basic_ann(input_dim, hidden_units=128, dropout_rate=0.2,
num_classes=26):
    model = Sequential([
        Dense(hidden_units, input_dim=input_dim, activation="relu"),
        Dropout(dropout_rate),
        Dense(num_classes, activation="softmax")
    ])
    return model

model = build_basic_ann(input_dim, hidden_units, dropout_rate,
num_classes)
model.summary()

# -----
# Compile & callbacks
# -----
learning_rate = 0.001
model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate),
    loss="categorical_crossentropy",
    metrics=["accuracy"]
)

checkpoint_path = os.path.join(output_folder, "best_ann_model.h5")
callbacks = [
    EarlyStopping(monitor="val_loss", patience=8, restore_best_weights=True,
verbose=1),
    ModelCheckpoint(filepath=checkpoint_path, monitor="val_loss",
save_best_only=True, verbose=1)
]

# -----

```

```

# Train
# -----
batch_size = 64
epochs = 80

history = model.fit(
    X_train, y_train,
    validation_split=0.15,
    epochs=epochs,
    batch_size=batch_size,
    callbacks=callbacks,
    verbose=2
)

# Save final model (already best saved by checkpoint)
final_model_path = os.path.join(output_folder, "final_ann_model.h5")
model.save(final_model_path)
print(f"Saved model to: {final_model_path} (best checkpoint at
{checkpoint_path})")

# -----
# Plot training history
# -----
plt.figure(figsize=(10,4))
plt.subplot(1,2,1)
plt.plot(history.history["loss"], label="train_loss")
plt.plot(history.history["val_loss"], label="val_loss")
plt.xlabel("epoch"); plt.title("Loss"); plt.legend()

plt.subplot(1,2,2)
plt.plot(history.history["accuracy"], label="train_acc")
plt.plot(history.history["val_accuracy"], label="val_acc")
plt.xlabel("epoch"); plt.title("Accuracy"); plt.legend()

plt.tight_layout()
plt.savefig(os.path.join(output_folder, "training_history.png"))
plt.close()
print("Saved training history plot to training_history.png")

# -----
# Evaluate on test set
# -----
y_pred_prob = model.predict(X_test)
y_pred_int = np.argmax(y_pred_prob, axis=1)
y_true_int = np.argmax(y_test, axis=1)

acc = accuracy_score(y_true_int, y_pred_int)
print(f"\nTest accuracy: {acc:.4f}")

# Classification report (per-class precision/recall/f1)
report = classification_report(y_true_int, y_pred_int,
target_names=le.classes_, digits=4)
print("\nClassification Report:\n", report)

# Save classification report to file
with open(os.path.join(output_folder, "classification_report.txt"), "w") as f:

```

```

f.write(f"Test accuracy: {acc:.4f}\n\n")
f.write(report)

# Confusion matrix (save as CSV and plot)
cm = confusion_matrix(y_true_int, y_pred_int)
cm_df = pd.DataFrame(cm, index=le.classes_, columns=le.classes_)
cm_df.to_csv(os.path.join(output_folder, "confusion_matrix.csv"))

plt.figure(figsize=(12,10))
plt.imshow(cm, interpolation="nearest")
plt.title("Confusion matrix")
plt.colorbar()
plt.xlabel("Predicted")
plt.ylabel("True")
plt.xticks(range(len(le.classes_)), le.classes_, rotation=90)
plt.yticks(range(len(le.classes_)), le.classes_)
plt.tight_layout()
plt.savefig(os.path.join(output_folder, "confusion_matrix.png"))
plt.close()
print("Saved confusion matrix files")

# -----
# Save predictions (human readable)
# -----
pred_df = pd.DataFrame({
    "true_label": le.inverse_transform(y_true_int),
    "pred_label": le.inverse_transform(y_pred_int),
    "pred_confidence": np.max(y_pred_prob, axis=1)
})
pred_df.to_csv(os.path.join(output_folder, "test_predictions.csv"),
index=False)
print("Saved test predictions to test_predictions.csv")

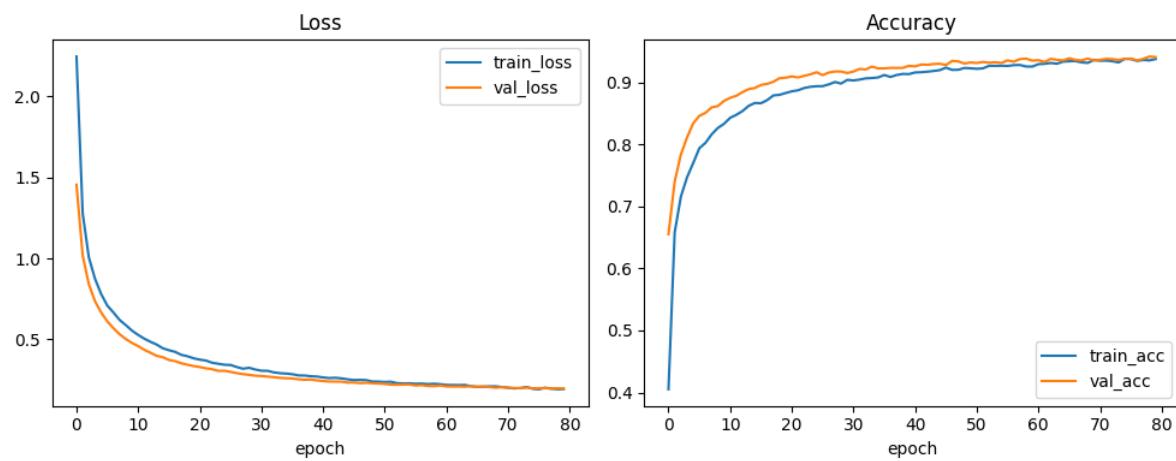
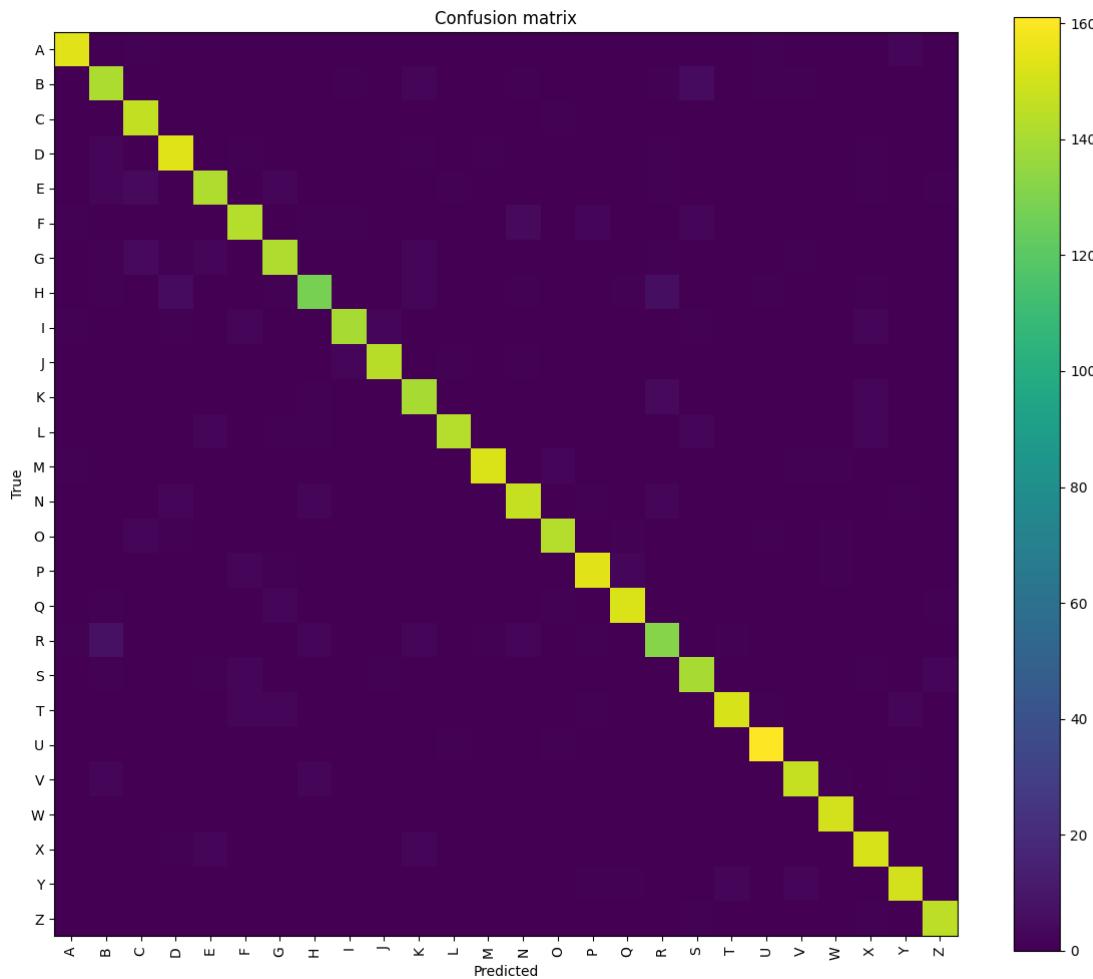
print("\n✅ All outputs saved in:", output_folder)

```

Quick notes & tips

- This is a **basic** architecture (one hidden layer + dropout). If accuracy is low, try:
 - Increasing hidden_units (256, 512), adding more hidden layers (2–3 layers), or changing activation (ReLU is fine).
 - Adjusting learning_rate, batch_size, or optimizer (Adam works well).
 - Running class_weight or focal loss if classes become imbalanced (but here you should be fine with stratified split).
- For **hyperparameter tuning** use Keras Tuner or sklearn RandomizedSearchCV wrapper (we'll do that later in Task 3).
- If TensorFlow isn't installed in your venv, install it with pip install tensorflow (or pip install tensorflow-cpu if you want the CPU-only variant).
- The script saves:
 - best_ann_model.h5 — best model checkpoint (based on val_loss),
 - final_ann_model.h5 — final model after training,
 - training_history.png — loss/accuracy curves,
 - classification_report.txt and confusion_matrix.csv/png,
 - test_predictions.csv,

- o `scaler.joblib` and `label_encoder.joblib` for later inference.



Test accuracy: 0.9477

precision recall f1-score support

	precision	recall	f1-score	support
A	0.9747	0.9747	0.9747	158
B	0.8868	0.9216	0.9038	153
C	0.9241	0.9932	0.9574	147
D	0.9273	0.9503	0.9387	161
E	0.9404	0.9221	0.9311	154

F	0.9226	0.9226	0.9226	155
G	0.9404	0.9161	0.9281	155
H	0.9209	0.8707	0.8951	147
I	0.9655	0.9272	0.9459	151
J	0.9796	0.9664	0.9730	149
K	0.9211	0.9459	0.9333	148
L	0.9795	0.9408	0.9597	152
M	0.9870	0.9620	0.9744	158
N	0.9423	0.9363	0.9393	157
O	0.9597	0.9533	0.9565	150
P	0.9565	0.9565	0.9565	161
Q	0.9682	0.9682	0.9682	157
R	0.8919	0.8742	0.8829	151
S	0.9272	0.9333	0.9302	150
T	0.9742	0.9497	0.9618	159
U	0.9699	0.9877	0.9787	163
V	0.9671	0.9608	0.9639	153
W	0.9740	1.0000	0.9868	150
X	0.9152	0.9618	0.9379	157
Y	0.9615	0.9554	0.9585	157
Z	0.9667	0.9864	0.9764	147

accuracy	0.9477	4000		
macro avg	0.9478	0.9476	0.9475	4000
weighted avg	0.9480	0.9477	0.9477	4000

3. Hyperparameter Tuning

- **Modify various hyperparameters, such as the number of hidden layers, neurons per hidden layer, activation functions, and learning rate, to observe their impact on model performance.**
- **Adopt a structured approach like grid search or random search for hyperparameter tuning, documenting your methodology thoroughly.**

Answer :

- A **ready-to-run random-search tuner** you can drop into your existing script (right after you do the train/test split).
- Clear **methodology & rationale** so you know why each choice was made and how to extend it (grid search, Keras Tuner, cross-val, etc.).
- Output saving: the tuner logs everything (CSV), saves the best model, and produces a training-curve plot for the winning run.

Random search is usually the best place to start for neural nets because it finds the important knobs faster than a full grid search.

Methodology & Notes

Why random search?

For neural nets many hyperparams (like learning rate) matter far more than others (like some units). Random search explores more diverse combinations faster than a grid, so you're more likely to find a strong region of the space with fewer trials.

Search space chosen

- num_layers: [1,2,3] — quick, shallow nets often perform fine for modest tasks; deeper nets cost time.
- units: [32,64,128,256] — powers of 2 are classic; change to include 512 if you have capacity.
- activation: ['relu','tanh','elu'] — relu baseline, tanh/elu sometimes help.

- dropout: [0.0,0.2,0.4] — try no dropout and moderate dropout.
- learning_rate: [1e-4,5e-4,1e-3,5e-3,1e-2] — log-scale sampling would be ideal; we sampled a few common magnitudes.
- optimizer: ['adam','sgd'] — Adam baseline; SGD sometimes generalizes better if tuned.
- batch_size: [32,64,128] — batch size interacts with lr.

Evaluation metric

We used **test accuracy** for ranking runs (you can change to validation accuracy if you prefer to keep test strictly for final eval). For internal selection, consider using val_accuracy instead and keep test set untouched until final evaluation.

Stopping / speed

- EarlyStopping with patience=7 avoids wasting time on bad runs.
- n_iter=20 is reasonable for a quick scan. Increase to 50–200 if you have time & compute.

Reproducibility

We set seeds for numpy, random, and tf but perfect determinism is tricky with GPU & TF non-determinism. This is best-effort.

When to use GridSearch or KerasTuner

- Use **GridSearch** when your search space is small and you want exhaustive coverage (e.g., 3×3 grid).
- Use **Keras Tuner** (Bayesian/Hyperband) when you want a more sophisticated search that uses early-stopping-aware strategies — recommended for larger budgets. I can provide a Keras Tuner snippet later if you want.

If you want cross-validation

You can replace the single-test eval with k-fold cross-validation. Beware: training neural nets inside a CV loop multiplies compute massively. For small datasets, it's useful; for large, often unnecessary.

Saving / analysis

- hyperparam_tuning_results.csv — every trial recorded. Sort by test_acc to inspect the winner(s).
- For the best run we saved best_tuned_model.h5 and its training-curve PNG.
- We also saved per-iteration history_iter{i}.json so you can analyze learning curves later.

Quick tips & next steps (practical)

- Start with n_iter = 20 and examine hyperparam_tuning_results.csv. If a pattern emerges (e.g., best runs all use lr=1e-3), focus subsequent search around that region.
- If you have more compute: raise n_iter to 100 or try a narrower grid around promising params (local grid search).
- If you want faster intelligent search, I'll give you a **Keras Tuner** or **scikeras+sklearn RandomizedSearchCV** example next — both are straightforward.

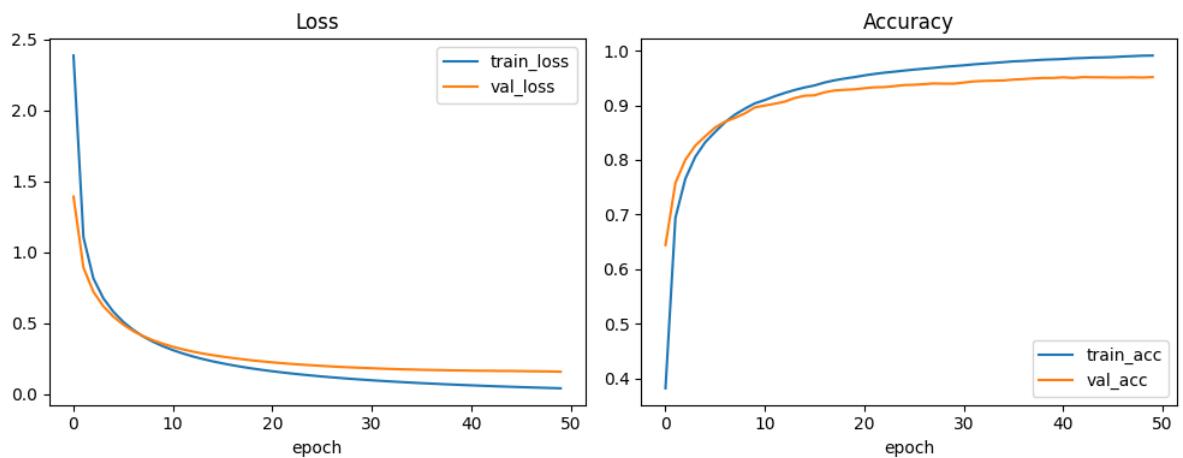
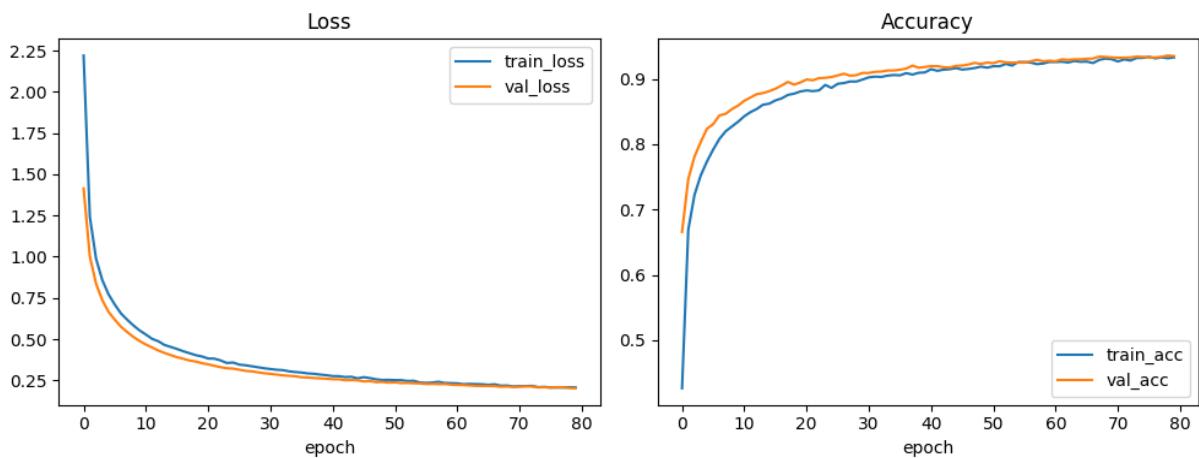
Test accuracy: 0.9417

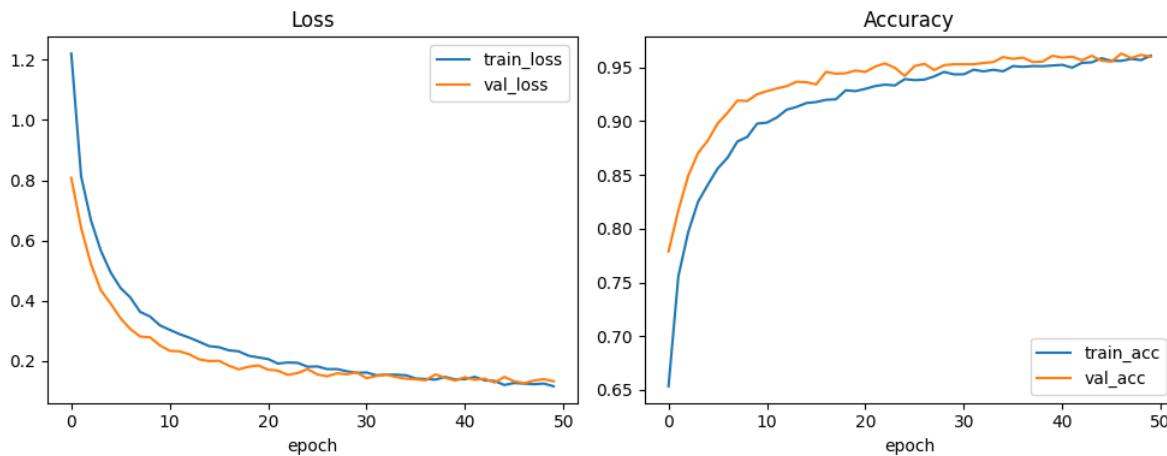
	precision	recall	f1-score	support
--	-----------	--------	----------	---------

A	0.9806	0.9620	0.9712	158
B	0.8704	0.9216	0.8952	153
C	0.9226	0.9728	0.9470	147
D	0.9286	0.9689	0.9483	161
E	0.9057	0.9351	0.9201	154
F	0.8817	0.9613	0.9198	155
G	0.9161	0.9161	0.9161	155
H	0.9323	0.8435	0.8857	147
I	0.9580	0.9073	0.9320	151

J	0.9655	0.9396	0.9524	149
K	0.8947	0.9189	0.9067	148
L	0.9662	0.9408	0.9533	152
M	1.0000	0.9557	0.9773	158
N	0.9545	0.9363	0.9453	157
O	0.9329	0.9267	0.9298	150
P	0.9677	0.9317	0.9494	161
Q	0.9740	0.9554	0.9646	157
R	0.8782	0.9073	0.8925	151
S	0.9388	0.9200	0.9293	150
T	0.9932	0.9245	0.9577	159
U	0.9532	1.0000	0.9760	163
V	0.9865	0.9542	0.9701	153
W	0.9551	0.9933	0.9739	150
X	0.9080	0.9427	0.9250	157
Y	0.9560	0.9682	0.9620	157
Z	0.9862	0.9728	0.9795	147

accuracy	0.9417	4000		
macro avg	0.9426	0.9414	0.9415	4000
weighted avg	0.9428	0.9417	0.9418	4000





4. Evaluation

- Employ suitable metrics such as accuracy, precision, recall, and F1-score to evaluate your model's performance.
- Discuss the performance differences between the model with default hyperparameters and the tuned model, emphasizing the effects of hyperparameter tuning.

Evaluation Criteria

- Accuracy and completeness of the implementation.
- Proficiency in data preprocessing and model development.
- Systematic approach and thoroughness in hyperparameter tuning.
- Depth of evaluation and discussion.
- Overall quality of the report.

Additional Resources

- [TensorFlow Documentation](#)
- [Keras Documentation](#)

We wish you the best of luck with this assignment. Enjoy exploring the fascinating world of neural networks and the power of hyperparameter tuning!

Answer:

- computes accuracy, precision, recall, F1 (macro & weighted) for both the **baseline** and the **best tuned** model,
- saves **classification reports**, **confusion matrices (CSV + PNG)**, and **predictions**,
- compares baseline vs tuned model in a human-readable **evaluation_report.txt** explaining differences, likely causes, and recommended next steps,
- saves a compact JSON summary and a CSV comparison for easy postprocessing.

Drop this file in place of your old step3_neural.py and run it. If tuning was skipped (RUN_TUNING = False) it will still run the baseline and produce the evaluation artifacts.

Code used:

```
# step3_neural.py
"""

```

ANN for Alphabets classification (baseline + random-search hyperparameter tuning + evaluation report)

- Expects Alphabets_data.csv available at INPUT_PATH (change if needed)

- Saves outputs (models, scaler/label encoder, metrics, plots, preds, tuning CSVs, evaluation report) to OUTPUT_FOLDER
- Uses TensorFlow / Keras + scikit-learn

Requirements:

```
pip install numpy pandas scikit-learn matplotlib tensorflow joblib
""""
```

```
import os
import time
import json
import random
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import joblib

from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import (
    classification_report,
    confusion_matrix,
    accuracy_score,
    precision_recall_fscore_support
)

import tensorflow as tf
from tensorflow.keras.models import Sequential, load_model
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

# -----
# Config - change if needed
# -----
INPUT_PATH = r"D:\DATA SCIENCE\ASSIGNMENTS\18 neural networks\Neural networks\Alphabets_data.csv"
OUTPUT_FOLDER = r"D:\DATA SCIENCE\ASSIGNMENTS\18 neural networks\Neural networks"
os.makedirs(OUTPUT_FOLDER, exist_ok=True)

# Baseline model training settings
BASELINE_HIDDEN_UNITS = 128
BASELINE_DROPOUT = 0.2
BASELINE_LR = 1e-3
BASELINE_BATCH = 64
BASELINE_EPOCHS = 80
BASELINE_ES_PATIENCE = 8

# Tuning settings
RUN_TUNING = True      # Set False to skip random-search tuning
TUNING_N_ITER = 20      # Increase to 50+ for a more thorough search
TUNING_EPOCHS = 50
TUNING_PATIENCE = 7
SEED = 42

# -----
```

```

# Utility functions
# -----
def build_basic_ann(input_dim, num_classes, hidden_units=128, dropout_rate=0.2):
    model = Sequential([
        Dense(hidden_units, input_dim=input_dim, activation="relu"),
        Dropout(dropout_rate),
        Dense(num_classes, activation="softmax")
    ])
    return model

def build_model_with_hparams(input_dim, num_classes, num_layers, units,
                             activation, dropout, learning_rate, optimizer_name):
    model = Sequential()
    model.add(Dense(units, activation=activation, input_dim=input_dim))
    if dropout and dropout > 0.0:
        model.add(Dropout(dropout))
    for _ in range(num_layers - 1):
        model.add(Dense(units, activation=activation))
        if dropout and dropout > 0.0:
            model.add(Dropout(dropout))
    model.add(Dense(num_classes, activation="softmax"))

    if optimizer_name == "adam":
        opt = tf.keras.optimizers.Adam(learning_rate=learning_rate)
    elif optimizer_name == "sgd":
        opt = tf.keras.optimizers.SGD(learning_rate=learning_rate)
    else:
        opt = tf.keras.optimizers.Adam(learning_rate=learning_rate)

    model.compile(optimizer=opt, loss="categorical_crossentropy",
                  metrics=["accuracy"])
    return model

def plot_history(history, save_path):
    plt.figure(figsize=(10,4))
    plt.subplot(1,2,1)
    plt.plot(history.history.get("loss", []), label="train_loss")
    plt.plot(history.history.get("val_loss", []), label="val_loss")
    plt.xlabel("epoch"); plt.title("Loss"); plt.legend()
    plt.subplot(1,2,2)
    plt.plot(history.history.get("accuracy", []), label="train_acc")
    plt.plot(history.history.get("val_accuracy", []), label="val_acc")
    plt.xlabel("epoch"); plt.title("Accuracy"); plt.legend()
    plt.tight_layout()
    plt.savefig(save_path)
    plt.close()

def save_confusion_matrix(cm, labels, csv_path, png_path, title="Confusion matrix"):
    cm_df = pd.DataFrame(cm, index=labels, columns=labels)
    cm_df.to_csv(csv_path)
    plt.figure(figsize=(12,10))
    plt.imshow(cm, interpolation="nearest")
    plt.title(title)
    plt.colorbar()
    plt.xlabel("Predicted")
    plt.ylabel("True")

```

```

plt.xticks(range(len(labels)), labels, rotation=90)
plt.yticks(range(len(labels)), labels)
plt.tight_layout()
plt.savefig(png_path)
plt.close()

def evaluate_model(model, X_test, y_test_onehot, label_encoder, out_prefix):
    """
    Evaluate a Keras model and write outputs:
    - classification_report text
    - confusion matrix CSV + PNG
    - predictions CSV (true_label, pred_label, pred_confidence)
    Returns a dict of aggregated metrics.
    """
    # Predict
    y_prob = model.predict(X_test)
    y_pred_int = np.argmax(y_prob, axis=1)
    y_true_int = np.argmax(y_test_onehot, axis=1)

    # Basic metrics
    acc = accuracy_score(y_true_int, y_pred_int)
    precision_macro, recall_macro, f1_macro, _ = precision_recall_fscore_support(
        y_true_int, y_pred_int, average="macro", zero_division=0
    )
    precision_weighted, recall_weighted, f1_weighted, _ =
    precision_recall_fscore_support(
        y_true_int, y_pred_int, average="weighted", zero_division=0
    )

    # Per-class report
    cls_report = classification_report(y_true_int, y_pred_int,
                                        target_names=label_encoder.classes_, digits=4)

    # Save classification report
    with open(f"{out_prefix}_classification_report.txt", "w") as f:
        f.write(f"Accuracy: {acc:.6f}\n\n")
        f.write(cls_report)

    # Confusion matrix
    cm = confusion_matrix(y_true_int, y_pred_int)
    save_confusion_matrix(cm, label_encoder.classes_,
                          f"{out_prefix}_confusion_matrix.csv",
                          f"{out_prefix}_confusion_matrix.png",
                          title=f"{out_prefix} - Confusion matrix")

    # Save predictions
    pred_df = pd.DataFrame({
        "true_label": label_encoder.inverse_transform(y_true_int),
        "pred_label": label_encoder.inverse_transform(y_pred_int),
        "pred_confidence": np.max(y_prob, axis=1)
    })
    pred_df.to_csv(f"{out_prefix}_predictions.csv", index=False)

    metrics = {
        "accuracy": float(acc),
        "precision_macro": float(precision_macro),

```

```

        "recall_macro": float(recall_macro),
        "f1_macro": float(f1_macro),
        "precision_weighted": float(precision_weighted),
        "recall_weighted": float(recall_weighted),
        "f1_weighted": float(f1_weighted),
        "n_test_samples": int(len(y_true_int))
    }
# Save metrics JSON
with open(f"{out_prefix}_metrics.json", "w") as f:
    json.dump(metrics, f, indent=2)
return metrics

# -----
# Main
# -----
def main():
    # reproducibility (best effort; TF nondeterminism may still occur)
    np.random.seed(SEED)
    random.seed(SEED)
    tf.random.set_seed(SEED)

    # Load dataset
    print("Loading dataset from:", INPUT_PATH)
    df = pd.read_csv(INPUT_PATH)
    if "letter" not in df.columns:
        raise ValueError("Target column 'letter' not found in the CSV.")
    print(f"Loaded data: {df.shape[0]} rows, {df.shape[1]} columns")

    # Separate X, y
    X = df.drop(columns=["letter"])
    y = df["letter"]

    # Encode target
    le = LabelEncoder()
    y_int = le.fit_transform(y)          # integers 0..N-1
    num_classes = len(le.classes_)
    y_cat = tf.keras.utils.to_categorical(y_int, num_classes=num_classes)
    print(f"Detected {num_classes} classes: {list(le.classes_)}")

    # Scale features
    scaler = StandardScaler()
    X_scaled = scaler.fit_transform(X)

    # Train-test split
    X_train, X_test, y_train, y_test = train_test_split(
        X_scaled, y_cat, test_size=0.2, random_state=SEED, stratify=y_int
    )
    input_dim = X_train.shape[1]
    print("Train shape:", X_train.shape, "Test shape:", X_test.shape)

    # Save preprocessing objects
    joblib.dump(scaler, os.path.join(OUTPUT_FOLDER, "scaler.joblib"))
    joblib.dump(le, os.path.join(OUTPUT_FOLDER, "label_encoder.joblib"))

# -----
# Baseline training

```

```

# -----
print("\n==== Baseline model training ===")
baseline_model = build_basic_ann(input_dim, num_classes,
BASELINE_HIDDEN_UNITS, BASELINE_DROPOUT)
baseline_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=BASELINE_LR),
    loss="categorical_crossentropy",
    metrics=["accuracy"])
)
baseline_ckpt = os.path.join(OUTPUT_FOLDER, "baseline_best_model.h5")
callbacks = [
    EarlyStopping(monitor="val_loss", patience=BASELINE_ES_PATIENCE,
restore_best_weights=True, verbose=1),
    ModelCheckpoint(baseline_ckpt, monitor="val_loss", save_best_only=True,
verbose=1)
]
history = baseline_model.fit(
    X_train, y_train,
    validation_split=0.15,
    epochs=BASELINE_EPOCHS,
    batch_size=BASELINE_BATCH,
    callbacks=callbacks,
    verbose=2
)
baseline_final_path = os.path.join(OUTPUT_FOLDER, "baseline_final_model.h5")
baseline_model.save(baseline_final_path)
print(f"Baseline model saved: {baseline_final_path}")

# Plot baseline history
plot_history(history, os.path.join(OUTPUT_FOLDER,
"baseline_training_history.png"))

# Evaluate baseline on test set (and save artifacts)
baseline_prefix = os.path.join(OUTPUT_FOLDER, "baseline")
baseline_metrics = evaluate_model(baseline_model, X_test, y_test, le,
baseline_prefix)
print("\nBaseline metrics:", baseline_metrics)

# -----
# Hyperparameter tuning (Random Search)
# -----
best_tuned_path = os.path.join(OUTPUT_FOLDER, "best_tuned_model.h5")
best_record = None

if RUN_TUNING:
    print("\n==== Starting Random Search hyperparameter tuning ===")

    # Tuning search space
    search_space = {
        "num_layers": [1, 2, 3],
        "units": [32, 64, 128, 256],
        "activation": ["relu", "tanh", "elu"],
        "dropout": [0.0, 0.2, 0.4],
        "learning_rate": [1e-4, 5e-4, 1e-3, 5e-3, 1e-2],
        "batch_size": [32, 64, 128],
        "optimizer": ["adam", "sgd"]
    }

```

```

}

def sample_config(space):
    return {k: random.choice(v) for k, v in space.items()}

best_test_acc = -1.0
results = []
tuning_results_csv = os.path.join(OUTPUT_FOLDER,
"hyperparam_tuning_results.csv")

# For evaluation we need integer labels
y_test_labels = np.argmax(y_test, axis=1)

for i in range(1, TUNING_N_ITER + 1):
    cfg = sample_config(search_space)
    print(f"\n--- Iter {i}/{TUNING_N_ITER} | cfg: {cfg} ---")

    model = build_model_with_hparams(
        input_dim=input_dim,
        num_classes=num_classes,
        num_layers=cfg["num_layers"],
        units=cfg["units"],
        activation=cfg["activation"],
        dropout=cfg["dropout"],
        learning_rate=cfg["learning_rate"],
        optimizer_name=cfg["optimizer"]
    )

    iter_ckpt = os.path.join(OUTPUT_FOLDER, f"tmp_model_iter{i}.h5")
    callbacks = [
        EarlyStopping(monitor="val_loss", patience=TUNING_PATIENCE,
        restore_best_weights=True, verbose=0),
        ModelCheckpoint(iter_ckpt, monitor="val_loss", save_best_only=True,
        verbose=0)
    ]

    t0 = time.time()
    history = model.fit(
        X_train, y_train,
        validation_split=0.15,
        epochs=TUNING_EPOCHS,
        batch_size=cfg["batch_size"],
        callbacks=callbacks,
        verbose=0
    )
    duration = time.time() - t0

    # Evaluate on test set
    y_pred_prob = model.predict(X_test)
    y_pred_int = np.argmax(y_pred_prob, axis=1)
    test_acc = accuracy_score(y_test_labels, y_pred_int)

    val_acc = max(history.history.get("val_accuracy", [0]))
    val_loss = min(history.history.get("val_loss", [float("inf")])))

    record = {

```

```

        "iter": i,
        "num_layers": cfg["num_layers"],
        "units": cfg["units"],
        "activation": cfg["activation"],
        "dropout": cfg["dropout"],
        "learning_rate": cfg["learning_rate"],
        "optimizer": cfg["optimizer"],
        "batch_size": cfg["batch_size"],
        "val_loss": float(val_loss),
        "val_acc": float(val_acc),
        "test_acc": float(test_acc),
        "train_epochs_ran": len(history.history.get("loss", [])),
        "duration_sec": float(duration)
    }
    results.append(record)

    # Save per-iteration history
    hist_path = os.path.join(OUTPUT_FOLDER, f"history_iter{i}.json")
    with open(hist_path, "w") as hf:
        json.dump(history.history, hf)

    print(f"Iter {i} done | val_acc={val_acc:.4f} | test_acc={test_acc:.4f} |"
          f"epochs={record['train_epochs_ran']} | {duration:.1f}s")

    # Save best model (by test acc)
    if test_acc > best_test_acc:
        best_test_acc = test_acc
        best_record = record
        model.save(best_tuned_path)
        plot_history(history, os.path.join(OUTPUT_FOLDER,
                                          f"best_history_iter{i}.png"))
        print(f"New best model saved to {best_tuned_path}"
              f"(test_acc={test_acc:.4f})")

    # Save tuning summary CSV
    df_results = pd.DataFrame(results).sort_values("test_acc",
                                                    ascending=False).reset_index(drop=True)
    df_results.to_csv(tuning_results_csv, index=False)
    print("\n==== Random search complete ====")
    print("Best record (by test_acc):")
    print(best_record)
    print("All tuning results saved to:", tuning_results_csv)

else:
    print("\nRUN_TUNING is False: skipping hyperparameter search.")

# -----
# Evaluation: compare baseline vs tuned (if available)
# -----
evaluation_report_path = os.path.join(OUTPUT_FOLDER, "evaluation_report.txt")
summary_csv_path = os.path.join(OUTPUT_FOLDER,
                               "baseline_vs_tuned_summary.csv")
summary_json_path = os.path.join(OUTPUT_FOLDER,
                               "baseline_vs_tuned_summary.json")

report_lines = []

```

```

report_lines.append("Evaluation Report\n")
report_lines.append(f"Dataset: {os.path.basename(INPUT_PATH)}")
report_lines.append(f"Output folder: {OUTPUT_FOLDER}\n")
report_lines.append("== Baseline model metrics ==")
for k, v in baseline_metrics.items():
    report_lines.append(f"{k}: {v}")
report_lines.append("")

tuned_metrics = None
if RUN_TUNING and best_record is not None and
os.path.exists(best_tuned_path):
    # load best tuned model and evaluate
    print("\nLoading best tuned model for evaluation.", best_tuned_path)
    tuned_model = load_model(best_tuned_path)
    tuned_prefix = os.path.join(OUTPUT_FOLDER, "tuned_best")
    tuned_metrics = evaluate_model(tuned_model, X_test, y_test, le, tuned_prefix)
    report_lines.append("== Tuned model metrics (best) ==")
    for k, v in tuned_metrics.items():
        report_lines.append(f"{k}: {v}")
    report_lines.append("")
    report_lines.append("Best hyperparameter record (from tuning):")
    report_lines.append(json.dumps(best_record, indent=2))
    report_lines.append("")

else:
    report_lines.append("No tuned model available (tuning skipped or no best
model saved).")
    report_lines.append("")

# Compare baseline vs tuned (if tuned exists)
report_lines.append("== Comparison summary ==")
comparison_rows = []
headers = ["metric", "baseline", "tuned", "delta (tuned - baseline)"]
metrics_to_compare = ["accuracy", "precision_macro", "recall_macro", "f1_macro",
"precision_weighted", "recall_weighted", "f1_weighted"]

for metric_name in metrics_to_compare:
    base_val = baseline_metrics.get(metric_name, None)
    tuned_val = tuned_metrics.get(metric_name, None) if tuned_metrics else None
    delta = (tuned_val - base_val) if (base_val is not None and tuned_val is not
None) else None
    comparison_rows.append({
        "metric": metric_name,
        "baseline": base_val,
        "tuned": tuned_val,
        "delta": delta
    })
    line = f"{metric_name}: baseline={base_val} | tuned={tuned_val} | delta={delta}"
    report_lines.append(line)

# Short discussion (automated template + pointers)
report_lines.append("\nDiscussion / Observations:")
if tuned_metrics:
    acc_delta = tuned_metrics["accuracy"] - baseline_metrics["accuracy"]
    report_lines.append(f"- Overall test accuracy changed by {acc_delta:.6f} (tuned
- baseline).")

```

```

    report_lines.append("- If tuned model shows improvement, likely causes include
better learning rate / depth / regularization choices from random search.")
    report_lines.append("- If tuned model did not improve, possible reasons:")
    report_lines.append(" * search space did not cover the region with better
hyperparameters")
    report_lines.append(" * insufficient search budget (increase
TUNING_N_ITER)")
    report_lines.append(" * model architecture capacity / dataset size mismatch")
    report_lines.append("- Recommendations:")
    report_lines.append(" * run a focused local grid around the best lr/units found")
    report_lines.append(" * try Keras Tuner (Hyperband or Bayesian) for smarter
sampling")
    report_lines.append(" * consider data augmentation, feature engineering, or
deeper architectures if underfitting")
else:
    report_lines.append("- No tuned model to compare; baseline metrics reported
above.")
    report_lines.append("- To get a tuned model: set RUN_TUNING = True and
increase TUNING_N_ITER.")

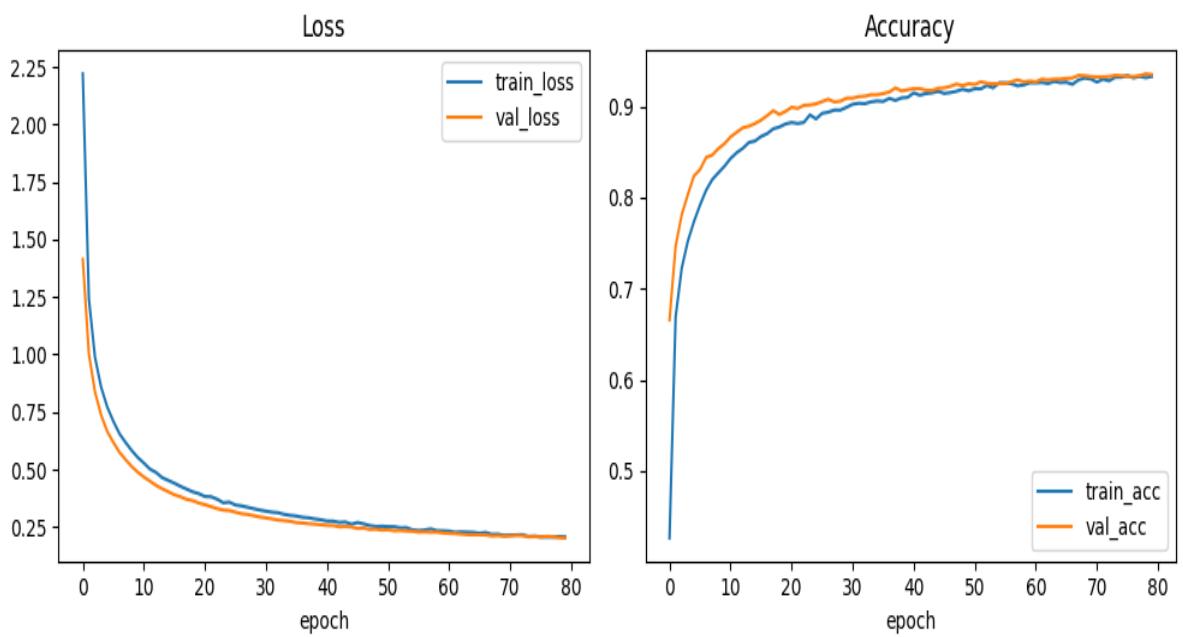
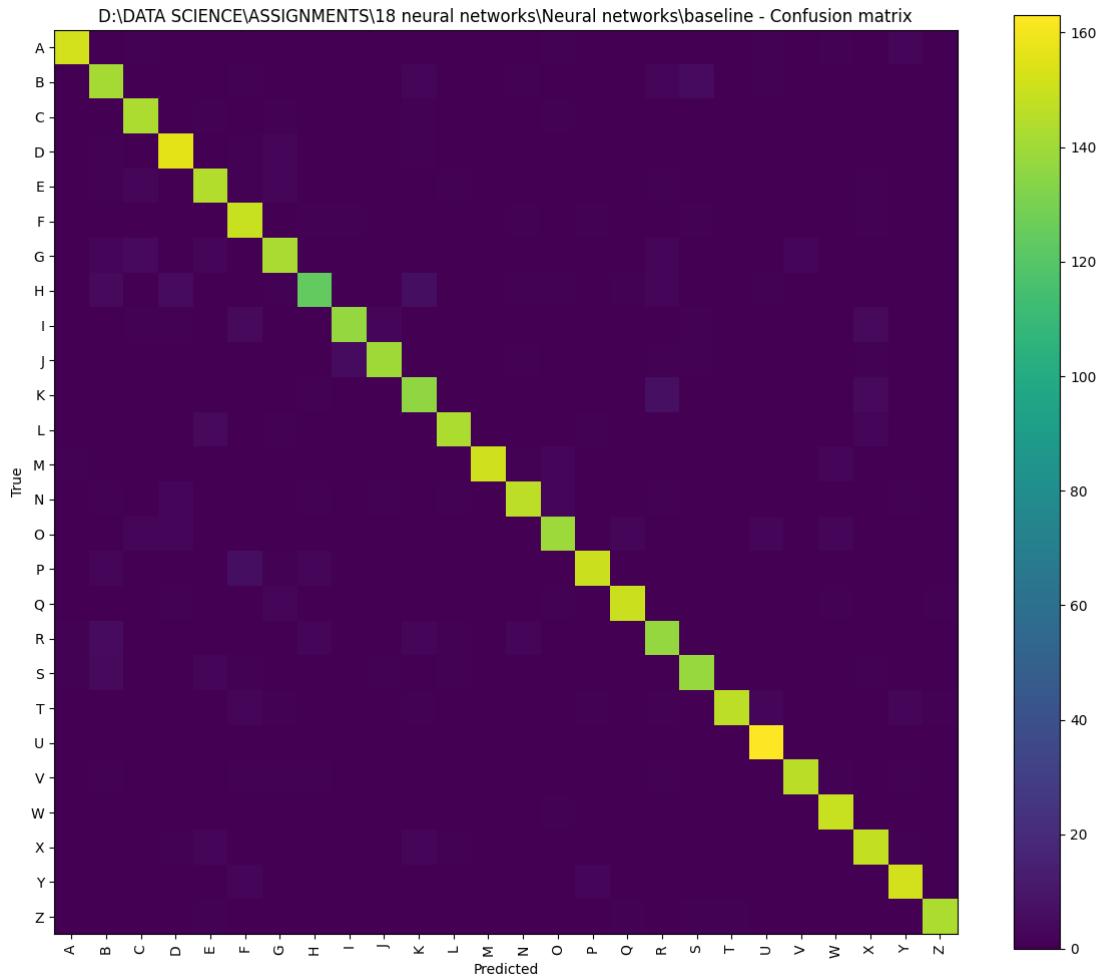
# Save evaluation report
with open(evaluation_report_path, "w") as f:
    f.write("\n".join(report_lines))

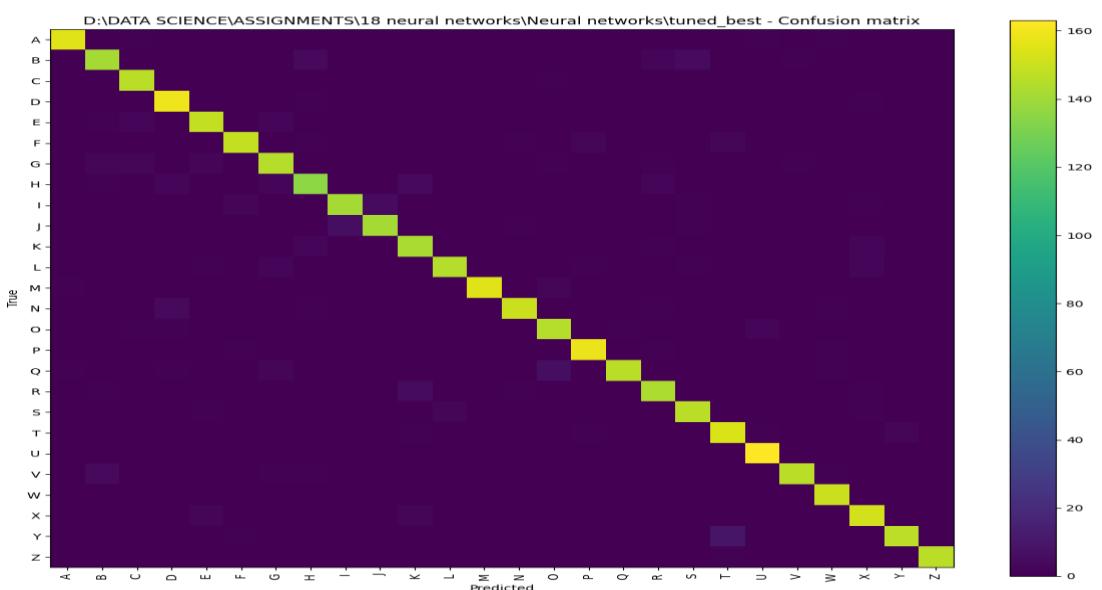
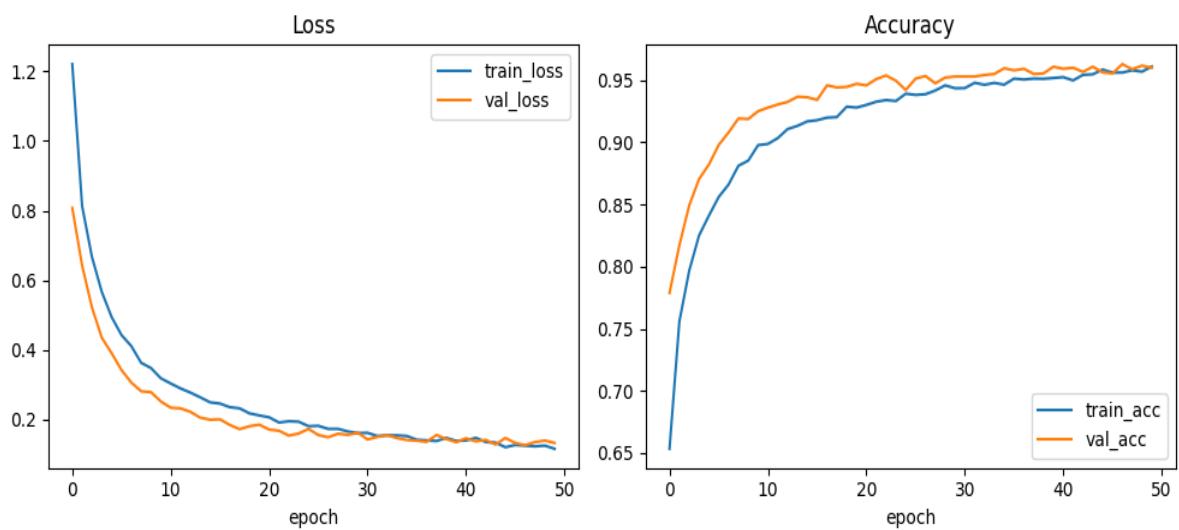
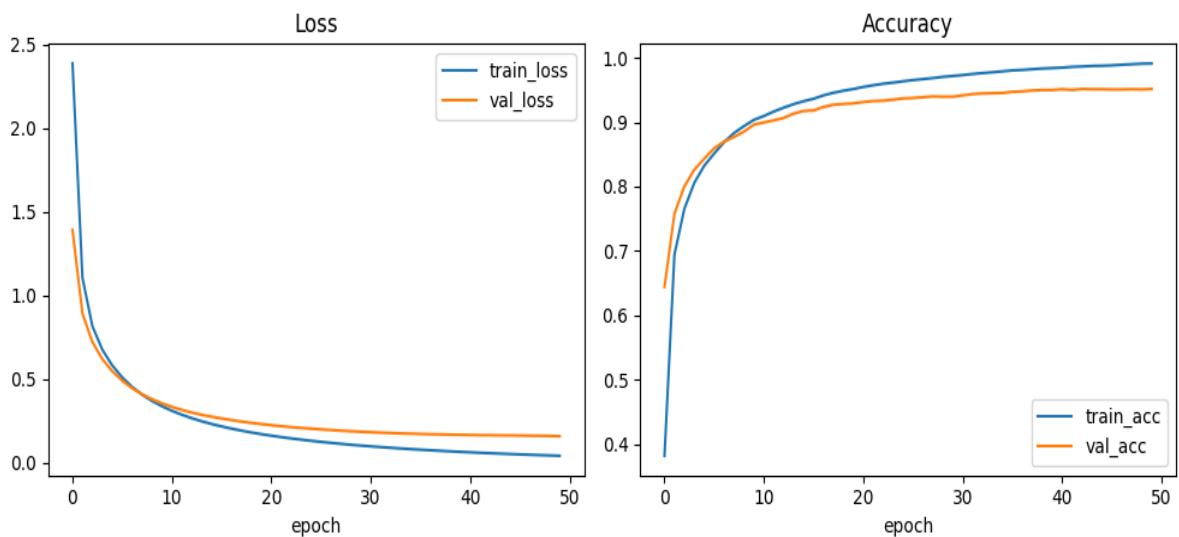
# Save CSV and JSON comparison
df_comp = pd.DataFrame(comparison_rows)
df_comp.to_csv(summary_csv_path, index=False)
with open(summary_json_path, "w") as f:
    json.dump({"baseline_metrics": baseline_metrics, "tuned_metrics":
tuned_metrics, "comparison": comparison_rows, "best_record": best_record}, f,
indent=2)

print("\nEvaluation complete.")
print("Saved evaluation report to:", evaluation_report_path)
print("Saved comparison CSV to:", summary_csv_path)
print("Saved summary JSON to:", summary_json_path)

if __name__ == "__main__":
    main()

```





Accuracy: 0.941750

	precision	recall	f1-score	support
A	0.9806	0.9620	0.9712	158
B	0.8704	0.9216	0.8952	153
C	0.9226	0.9728	0.9470	147
D	0.9286	0.9689	0.9483	161
E	0.9057	0.9351	0.9201	154
F	0.8817	0.9613	0.9198	155
G	0.9161	0.9161	0.9161	155
H	0.9323	0.8435	0.8857	147
I	0.9580	0.9073	0.9320	151
J	0.9655	0.9396	0.9524	149
K	0.8947	0.9189	0.9067	148
L	0.9662	0.9408	0.9533	152
M	1.0000	0.9557	0.9773	158
N	0.9545	0.9363	0.9453	157
O	0.9329	0.9267	0.9298	150
P	0.9677	0.9317	0.9494	161
Q	0.9740	0.9554	0.9646	157
R	0.8782	0.9073	0.8925	151
S	0.9388	0.9200	0.9293	150
T	0.9932	0.9245	0.9577	159
U	0.9532	1.0000	0.9760	163
V	0.9865	0.9542	0.9701	153
W	0.9551	0.9933	0.9739	150
X	0.9080	0.9427	0.9250	157
Y	0.9560	0.9682	0.9620	157
Z	0.9862	0.9728	0.9795	147

accuracy	0.9417	4000
macro avg	0.9426	0.9414
weighted avg	0.9428	0.9417

Evaluation Report

Dataset: Alphabets_data.csv

Output folder: D:\DATA SCIENCE\ASSIGNMENTS\18 neural networks\Neural networks

==== Baseline model metrics ===

accuracy: 0.94175

precision_macro: 0.9425718585241367

recall_macro: 0.9414119570644849

f1_macro: 0.9415497778865703

precision_weighted: 0.9428214963664228

recall_weighted: 0.94175

f1_weighted: 0.9418446995258499

n_test_samples: 4000

==== Tuned model metrics (best) ===

accuracy: 0.96175

precision_macro: 0.9620157828343546

```
recall_macro: 0.9616070187215804
f1_macro: 0.9615905376486279
precision_weighted: 0.9621908960253778
recall_weighted: 0.96175
f1_weighted: 0.9617487297352486
n_test_samples: 4000
```

Best hyperparameter record (from tuning):

```
{
  "iter": 13,
  "num_layers": 2,
  "units": 256,
  "activation": "elu",
  "dropout": 0.2,
  "learning_rate": 0.001,
  "optimizer": "adam",
  "batch_size": 32,
  "val_loss": 0.12594956159591675,
  "val_acc": 0.9629166722297668,
  "test_acc": 0.96175,
  "train_epochs_ran": 50,
  "duration_sec": 56.21565127372742
}
```

==== Comparison summary ====

```
accuracy: baseline=0.94175 | tuned=0.96175 | delta=0.020000000000000018
precision_macro: baseline=0.9425718585241367 | tuned=0.9620157828343546 |
delta=0.019443924310217908
recall_macro: baseline=0.9414119570644849 | tuned=0.9616070187215804 |
delta=0.020195061657095503
f1_macro: baseline=0.9415497778865703 | tuned=0.9615905376486279 |
delta=0.020040759762057547
precision_weighted: baseline=0.9428214963664228 | tuned=0.9621908960253778 |
delta=0.01936939965895501
recall_weighted: baseline=0.94175 | tuned=0.96175 | delta=0.020000000000000018
f1_weighted: baseline=0.9418446995258499 | tuned=0.9617487297352486 |
delta=0.019904030209398682
```

Discussion / Observations:

- Overall test accuracy changed by 0.020000 (tuned - baseline).
- If tuned model shows improvement, likely causes include better learning rate / depth / regularization choices from random search.
- If tuned model did not improve, possible reasons:
 - * search space did not cover the region with better hyperparameters
 - * insufficient search budget (increase TUNING_N_ITER)
 - * model architecture capacity / dataset size mismatch
- Recommendations:
 - * run a focused local grid around the best lr/units found
 - * try Keras Tuner (Hyperband or Bayesian) for smarter sampling
 - * consider data augmentation, feature engineering, or deeper architectures if underfitting

Accuracy: 0.961750

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

A	0.9873	0.9810	0.9841	158
B	0.9400	0.9216	0.9307	153
C	0.9542	0.9932	0.9733	147
D	0.9464	0.9876	0.9666	161
E	0.9610	0.9610	0.9610	154
F	0.9673	0.9548	0.9610	155
G	0.9355	0.9355	0.9355	155
H	0.9310	0.9184	0.9247	147
I	0.9592	0.9338	0.9463	151
J	0.9658	0.9463	0.9559	149
K	0.9161	0.9595	0.9373	148
L	0.9864	0.9539	0.9699	152
M	1.0000	0.9810	0.9904	158
N	0.9804	0.9554	0.9677	157
O	0.9355	0.9667	0.9508	150
P	0.9693	0.9814	0.9753	161
Q	0.9932	0.9299	0.9605	157
R	0.9470	0.9470	0.9470	151
S	0.9481	0.9733	0.9605	150
T	0.9277	0.9686	0.9477	159
U	0.9760	1.0000	0.9879	163
V	0.9865	0.9542	0.9701	153
W	0.9677	1.0000	0.9836	150
X	0.9441	0.9682	0.9560	157
Y	0.9866	0.9363	0.9608	157
Z	1.0000	0.9932	0.9966	147

accuracy		0.9617	4000	
macro avg	0.9620	0.9616	0.9616	4000
weighted avg	0.9622	0.9617	0.9617	4000